

Project 4

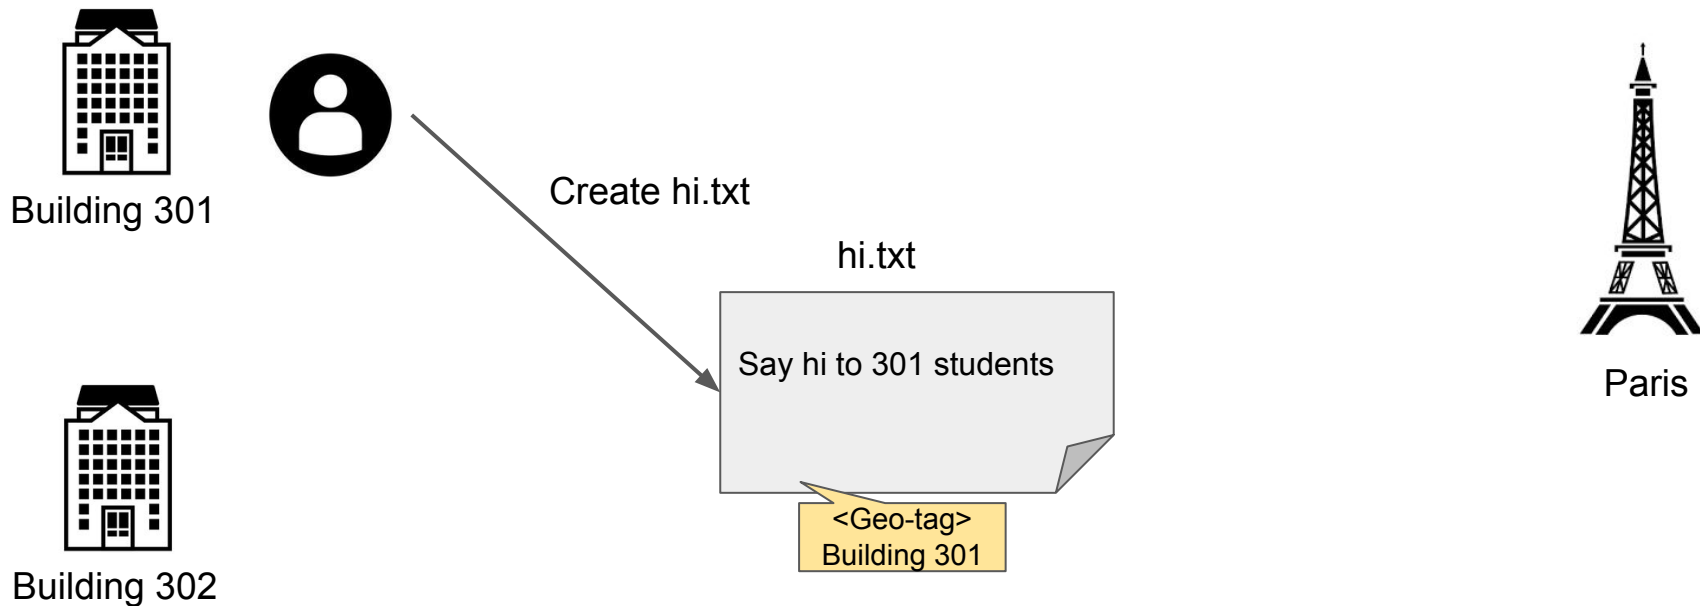
Hello, File System!

May 25, 2021
SNU Operating Systems

Overview

- Geo-tagged file system (based on only ext2)
 - Attach a GPS tag to each **regular file**
 - GPS tag should be set whenever a file is **created / modified**
- Access control with the GPS location
 - Files are **only accessible** from the location where they were most recently created/modified

Overview



Assume that 301 & 302 are
geometrically close

Overview



Building 301



Building 302



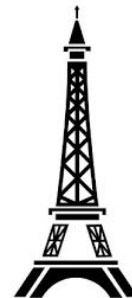
Append hi.txt

hi.txt

Say hi to 301 students
Say hi to 302 students

<Geo-tag>
Building 302

GPS location modified!



Paris

Assume that 301 & 302 are
geometrically close

Overview

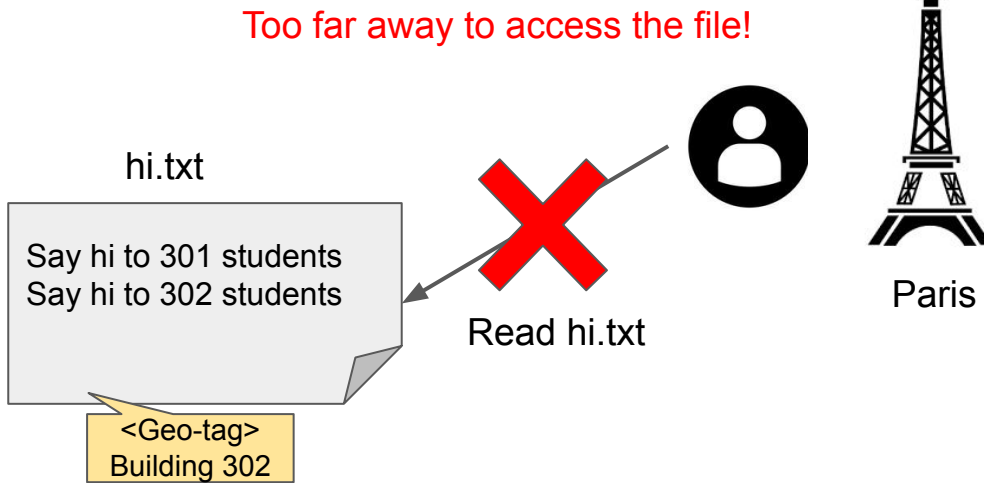


Building 301



Building 302

Assume that 301 & 302 are
geometrically close



☢ The `get_gps_location` syscall succeeds
regardless of the current location

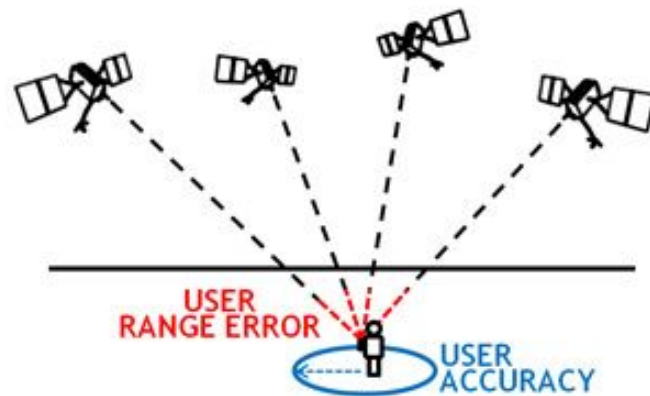
Four Steps of Project 4

- ① Syscall to set the device's current GPS location
- ② Add geo-tag support to the Linux inode structure
- ③ User-space testing for step ②
- ④ Access control based on GPS location

1. Set GPS location

- Struct `gps_location` has five fields
 - Latitude, Longitude, Accuracy
 - *_fractional accurate to sixth decimal places
- Implement system call `sys_set_gps_location`
 - Set current location of the device in the kernel
 - Recall what you did in project 2!

```
struct gps_location {  
    int lat_integer;  
    int lat_fractional;  
    int lng_integer;  
    int lng_fractional;  
    int accuracy;  
};
```



2. Add geo-tag to inode structure

- What to do with inode?
 - Whenever a file is created/modified in our custom ext2 file system, the current GPS location in the kernel should be recorded in the inode.
- What is “inode”?
 - Data structure that manages file metadata
 - Owner, date/time, permission control, file type, ...
 - Related operations are defined internally as `inode_operations`
- Its implementation in the kernel is very similar with `struct sched_class`

Implementation of sched_class

```
// kernel/sched/rt.c
const struct sched_class rt_sched_class = {
    .next = &fair_sched_class,
    .enqueue_task = enqueue_task_rt,
    .dequeue_task = dequeue_task_rt,
    .yield_task = yield_task_rt,

    .check_preempt_curr = check_preempt_curr_rt,

    .pick_next_task = pick_next_task_rt,
    .put_prev_task = put_prev_task_rt,

#ifdef CONFIG_SMP
    .select_task_rq = select_task_rq_rt,

    .set_cpus_allowed = set_cpus_allowed_rt,
    .rq_online = rq_online_rt,
    .rq_offline = rq_offline_rt,
    .pre_schedule = pre_schedule_rt,
    .post_schedule = post_schedule_rt,
    .task_woken = task_woken_rt,
    .switched_from = switched_from_rt,

    .set_curr_task = set_curr_task_rt,
    .task_tick = task_tick_rt,

    .get_rr_interval = get_rr_interval_rt,

    .prio_changed = prio_changed_rt,
    .switched_to = switched_to_rt,
#endif
};
```

Interface

Implementation

Multiple implementation sets

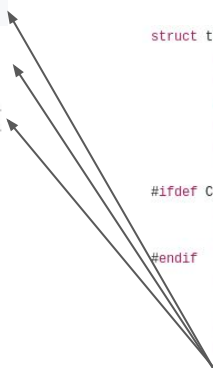
```
const struct sched_class rt_sched_class = {
const struct sched_class fair_sched_class = {
const struct sched_class idle_sched_class = {
```

Pointing from each task_struct

```
struct task_struct {
    volatile long state; /* -1 unrunnable
    void *stack;
    atomic_t usage;
    unsigned int flags; /* per process
    unsigned int ptrace;

#ifdef CONFIG_SMP
    struct llist_node wake_entry;
    int on_cpu;
#endif
    int on_rq;

    int prio, static_prio, normal_prio;
    unsigned int rt_priority;
    const struct sched_class *sched_class;
    struct sched_entity se;
    struct sched_rt_entity rt;
```



Similar for inode_operations

```
const struct inode_operations ext4_file_inode_operations = {
```

Interface

```
.setattr  
.getattr  
.setxattr  
.getxattr  
.listxattr  
.removexattr  
.get_acl  
.fiemap
```

Implementation

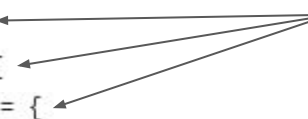
```
= ext4_setattr,  
= ext4_getattr,  
= generic_setxattr,  
= generic_getxattr,  
= ext4_listxattr,  
= generic_removexattr,  
= ext4_get_acl,  
= ext4_fiemap,
```

```
};
```

```
struct inode {  
    umode_t            i_mode;  
    unsigned short     i_opflags;  
    kuid_t             i_uid;  
    kgid_t             i_gid;  
    unsigned int       i_flags;  
  
#ifdef CONFIG_FS_POSIX_ACL  
    struct posix_acl    *i_acl;  
    struct posix_acl    *i_default_acl;  
#endif
```

```
const struct inode_operations ext3_dir_inode_operations = {  
const struct inode_operations ext4_file_inode_operations = {  
const struct inode_operations ext4_special_inode_operations = {
```

```
const struct inode_operations *i_op;  
struct super_block *i_sb;
```



2. Add geo-tag to inode: Modify inode operations

- Add the following two function pointer fields to the `inode_operations` structure in `include/linux/fs.h`
 - `int (*set_gps_location)(struct inode *);`
 - `int (*get_gps_location)(struct inode *, struct gps_location *);`
- Implement & register those functions with ext2 file system `inode_operations`
- Call the *setter* properly whenever the file is created or modified

2. Update GPS at inode: Modify inode structure

- Modify two inode structures in `fs/ext2/ext2.h`
 - `ext2_inode`: on-disk representation
 - `ext2_inode_info`: in-memory representation
- You have to add 5 attributes to each structure
- Modify some functions that transform between those two structures
- Be careful with physical(on-disk) representation
 - You should consider endianness, order of the fields, ...

3. User-space testing

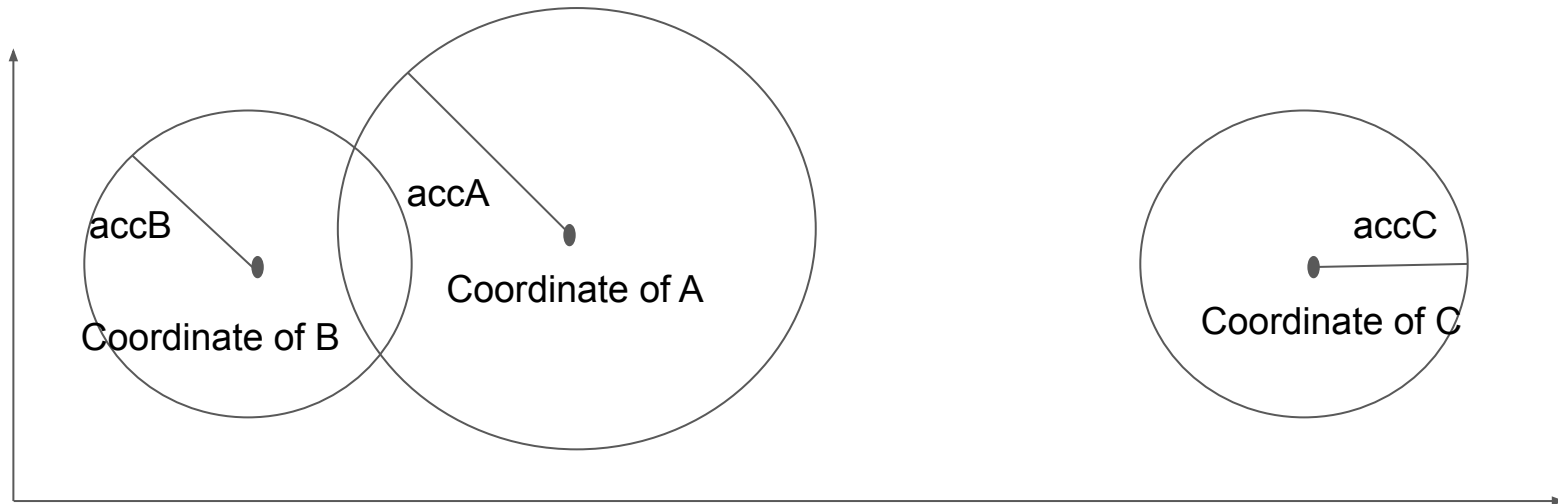
- Testing the GPS location setter/getter
 - By default, the root file system inside QEMU is btrfs
 - We will use e2fsprogs to test our custom ext2 file system!
- Create ext2 file system with mke2fs.
 - You need to modify the ext2 inode structure inside e2fsprogs to make mke2fs match your modified ext2 file system in the kernel
 - `e2fsprogs/lib/ext2fs/ext2_fs.h`

4. Location-based access control

- Files of the modified ext2 can be **only accessible** from the location close to where they are recently created/modified.
- When user tries to access to geo-tagged file,
 - Check based on existing access control mechanism first
 - If it passes the existing one, then check geometric proximity condition → this work as an extra condition!
- TIP: there is an ***inode operation used for general access control***. You can use it.

4. Location-based access control

- 'Accuracy' is used to determine whether the current location is *geometrically close* to file location
 - (A, B) is close
 - (A, C), (B, C) is not close



4. Location-based access control

- Remember: there are **no float/double types** in the kernel
- You should calculate the distance between two coordinates **with your own algorithm**.
- Document any assumptions or approximations on README.md

Be careful!

- The current device location is a shared mutable state, so you should properly synchronize access to it.
- Never access the user-space memory address directly.
 - Remind what we did in Project 1 :)
- For parameters in struct `gps_location` in the `set_gps_location` system call, make sure they are in their valid range.
 - Do NOT include 180.xxx for longitude!

About submission (IMPORTANT!)

- **No Design Review!**
- Three things to submit
 - `test/gpsupdate.c` : updates kernel GPS location with `sys_set_gps_locatin`
 - `test/file_loc.c`: get GPS location recorded in geo-tagged ext2 file
 - `proj4.fs` : should contain at least 1 directory & 2 files with different GPS coordinates

About submission (IMPORTANT!)

- Don't be late!
 - TAs will clone all repositories exactly at the deadline
- Submit code
 - Your team's private project 4 repo (swsnu/project4-hello-file-system-team-n), master branch
 - README: description of your implementation, how to build, and lessons learned
- Submit slides and demo (n is your team number)
 - Email: osspr2021@gmail.com
 - Title: [Project 4] Team n
 - Attachments: team-n-slides.{ppt,pdf}, team-n-demo.{mp4,avi,...}
 - One slide file, one demo video!

Q & A